

Using GPU as an accelerator

Anil Shanbhag

September 16, 2016

There is a lot of excitement around using accelerators like GPU and FPGAs to accelerate analytics. In fact, in the recent VLDB, two of the keynotes I attended (one by Ion and other by Gustavo) were centered around this theme.

GPUs have been widely adopted to do deep learning. However, GPUs also have a lot of potential to accelerate memory-bound applications like the ones we considered in NVL. There are two main contributing factors for why GPU's are appealing now:

1) The latest generation of GPU's have large amount of memory (the latest K80 Tesla card has 24GB memory, 3 years ago most papers had graphics card with at max 4 GB memory) and significantly higher memory bandwidth than a CPU (CPU bandwidth ~ 60 GBps compared to 480 GBps in K80). On a local machine with Titan X GPU, we observed a memory bandwidth of 280 GBps on GPU (listed as max 330 GBps in device spec) compared to 47GBps on CPU. The PCIe transfer speed is much slower. As a result, all the works so far, which had to ship data from CPU to GPU, suffered from limited gains as PCIe transfer time becomes the bottleneck. The large memory allows us to keep/cache all or a good fraction of the dataset on the GPU itself, eliminating the PCIe transfer overhead.

2) GPUs are becoming commodities. Azure recently launched the N-series which are machines with GPUs. The largest of them comes with 2 K80 cards having an aggregate of 48GB GPU memory. The price point is \$2.48ph for 24 cores/224GB RAM in addition to 2 GPUs. This is comparable to \$1.83 for 20 core/140GB RAM machine with no GPU. More cards can be stacked, MapD has machines with 8 K80 cards attached on the same machine.

There are a number of exciting research directions.

1 Generating Optimized Code

GPUs can theoretically accelerate analytics applications. In figure 1, we took the GPUDB database which was used in [2, 1] on the star schema

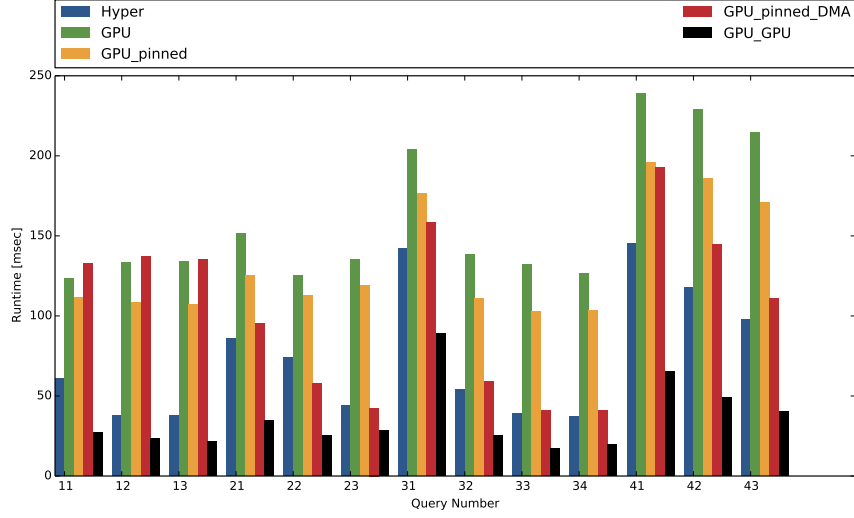


Figure 1: Comparing performance.

benchmark and compared its performance against Hyper. The paper compared the performance against MonetDB back in 2013 and observed at least 2x performance improvement with data stored in main memory. Much has changed since then, with improved algorithms for joins and query compilation pioneered by Hyper, main-memory databases are almost always memory-bound. The figure contains 4 modes in which GPUDB ran:

- GPU: data is on host and in paged memory
- GPU pinned: data is on host and in pinned memory
- GPU pinned DMA: data is on host, in pinned memory and data transfer done via DMA
- GPU GPU: data is on GPU

We observed that GPUDB never performs better than Hyper when data is in main memory and the gain is nowhere close to the bandwidth ratio when data is on GPU. In a hand-tuned implementation we observed greater than 5x gain for q11,q12,q13. There is a need to generate better code targetted for GPU. As machine learning applications already default to using GPUs, there is scope to do all-in GPU with cross library optimization.

2 Algorithms

There has been considerable research on optimizing standard relational operators for main-memory databases. The same is not true for GPUs. There

is still considerable room on coming up with better algorithms and deciding what is the algorithm to use for each operation. One operator I looked at so far is Top-k. On CPU, a sequential version can be done using a priority queue to have $O(n \log k)$. On the GPU, everyone just does a sort followed by selecting the top k entries. I have a sketch of an algorithm that has a $O(n(\log k)^2)$ work complexity and $\log k$ delay. There is probably some tricks to implement it efficiently and do come with similar algorithms for other operations.

3 Heterogenous Computing

Going back to the Azure VM, the machine has 20 cores, 224GB RAM and 2 K80 with 48GB RAM in total. Complete utilization of the machine requires using both CPU and GPU resources. There are a number of ways to do this. There are a number of ways to look at the problem too.

Scope

- Look at relational operations
- Look at general operators a.la tensorflow

Where is the data

- Data is split between GPU and CPU
- There are two copies of data: one on CPU, one on GPU
- Data is only on the CPU, moved to GPU

How to partition

- Horizontally partition the problem, process fraction of the data on GPU / fraction on CPU
- Look at inter-operator parallelism
- Look at query level parallelism: Query gets scheduled on GPU or CPU

The specific case I am looking at is relational operators with data being on the CPU. We will attempt to run the incoming query on the GPU and move/cache the columns used on the GPU. They are not deleted. When a subsequent query wants to run on the GPU, it will move any additional columns to the GPU in order to execute. Since many queries might access the same columns, we save on the transfer cost by caching it. When we exhaust the GPU memory, we can replace an existing column with the needed column. This is similar to caches in CPU land, but the difference is that there is one full copy of the data on the CPU. So, if the query requires

many new columns which we predict maybe by forecasting will not be used in future, we can execute the query on CPU and avoid moving out columns from the GPU.

There is no good open-source GPU-based database. The database resulting from this effort could be an useful resource. NVL could be a good substrate for compiling the query plans.

4 Cost Modelling

It turns out Tensorflow does not have any cost model. Atleast there is no such thing in the public release. User has to manually place computation on different devices. There has been some work on cost modelling of relational operators for GPU operators [2]. This was relatively easy due to the small number of operators. For arbitrary operators, it can possibly be done but would require much more work.

References

- [1] K. Wang, K. Zhang, Y. Yuan, S. Ma, R. Lee, X. Ding, and X. Zhang. Concurrent analytical query processing with gpus. *Proceedings of the VLDB Endowment*, 7(11):1011–1022, 2014.
- [2] Y. Yuan, R. Lee, and X. Zhang. The yin and yang of processing data warehousing queries on gpu devices. *Proceedings of the VLDB Endowment*, 6(10):817–828, 2013.